



Write-scalable, synchronous multi-master  
PostgreSQL cluster with shared nothing approach

Koichi Suzuki  
Mason Sharp



## Today's Talk

- What is Postgres-XC?
  - Concept and Ultimate Goal
- How to achieve read/write scalability
- Postgres-XC component
  - Global Transaction Manager
  - Coordinator
  - Data Node
- Current Status and Evaluation
- Possible Applications
- Issues and Roadmap



## What is Postgres-XC? (1)

- Write-scalable PostgreSQL cluster
  - More than 3 1/4 performance scalability with five servers, compared with pure PostgreSQL (DBT-1)
- Synchronous multi-master configuration
  - Any update to any master is visible from other masters immediately.
  - Not just a “replication”
    - Distribution (partition) and replication combination of tables



## What is Postgres-XC? (2)

- Table location transparent
  - Tables can be replicated or distributed (partitioned)
    - Best utilize parallelism among involved servers.
  - Can continue to use the same applications.
  - No change in transaction handling.
- Based upon PostgreSQL
- Same API to Apps. as PostgreSQL



## Postgres-XC Applications

- Short transaction applications (DBT-1/2 etc.)
  - Transactions can be executed in parallel in multiple data nodes.
- Complicated data warehouse (DBT-3 etc.)
  - Statement can be divided into several pieces executed in parallel in multiple data nodes.
    - (Statement handling not available yet.)




# Goal of Postgres-XC

Global, Transparent and  
Consistent View of Clustered  
Database

Read/Write  
Apps

Database  
Cluster





## Why Write-Scalability?

- Many application could be write-traffic bottleneck such as –
  - Access log in BLOG/SNS
  - Mission critical systems like internet shopping site, telephone customer billing, call information and securities trade
- Now application has to deal with such write bottleneck using multi-database.
  - Not distribution-transparent.
  - No cluster-specific application codes.
- As applications grow
  - It is desirable to make database distribution transparent for write operations too.
- But it is complicated ...



## Why Shared-Nothing?

- Most Cost-Efficient
  - Configurable only with commodity hardware
  - No need for dedicated disk system
  - No need for dedicated interface
- Flexibility to deploy
  - Can apply very simple to complicated cluster configuration





## Current Status and Plan

- Version 0.9 is available
  - <http://sourceforge.net/projects/postgres-xc/>
  - Simple statement w/o cross-node operation
    - Sufficient to run DBT-1 and pgbench
  - Create Table
- Version 0.91 – Now available (Wednesday)
  - Copy
  - Aggregate Functions
  - Replicated Table Update



## Current Status and Plan (cond.)

- Version 1.0 - July, 2010
  - Order By, Distinct
  - Subqueries
  - Views/Rules
  - DDLs
- Version 1.1 – Sept, 2010 (Still planning)
  - Installer
  - Operation Utilities
  - Dump/Restore (logical)
  - Cross-node operation (basic)
  - Extended Query Protocol (JDBC)
  - Global Timestamp
  - Simple Cursor (ECPG, JDBC, PHP, etc.)




## How to Achieve Read/Write Scalability

- Parallelism
  - Transactions run in parallel in database cluster
  - A statement can run in parallel in database cluster (future issue)
- Maintain Transaction Control
  - Transaction Timestamp (Transaction ID)
  - MVCC visibility
- Provide Global Values
  - Sequence
  - Timestamp (future issue)




# Parallel Transaction Execution





# Postgres-XC Configuration Outline





# Postgres-XC Components

- GTM (Global Transaction Manager)
  - Provide global transaction information to each transaction
    - Transaction ID
    - Snapshot
  - Provide other global data to statements
    - Sequence
    - Time/Sysdate (under plan)
- Coordinator
  - Parse statements and determine location of involved data
  - Transfer statements for each data node (if needed)
  - Application I/F
- Data Node
  - Store actual data
  - Execute statements from Coordinators



## Tables in Postgres-XC

- Tables are replicated or distributed
  - Replicated Table
    - Each Data Node stores whole replicated table.
    - Replication is maintained in the statement basis (not WAL basis)
  - Distributed Table
    - Each tuple is assigned a Data Node to go
      - Based on a value of a column (distribution key)
        - » Hash
        - » Round-Robin
        - » Range (future)
        - » User-Defined (future)



## How to Determine Distributed/Replicated?

- Transaction tables may be partitioned so that each transaction can be executed in limited number of data nodes.
- Master tables may be replicated so that each transaction can read row values locally.



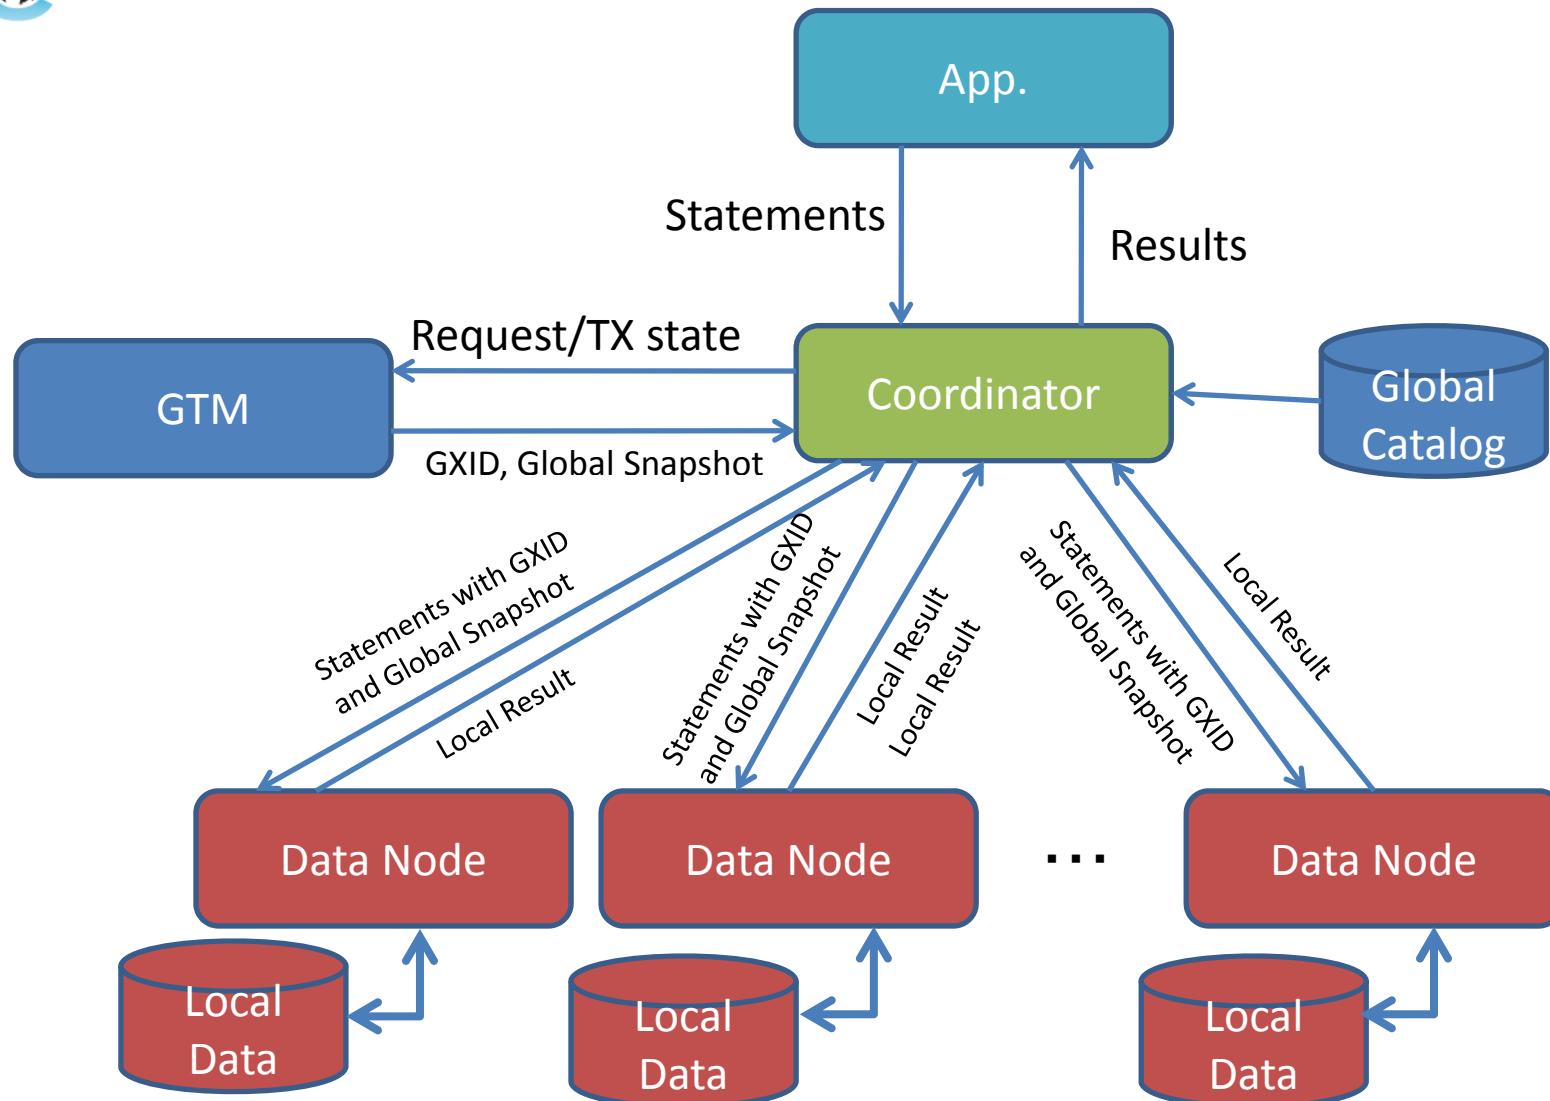


## GTM – Global Transaction Manager

- GTM is the key of Postgres-XC transaction management
  - Extracted essential of transaction management feature of PostgreSQL
    - Unique Transaction ID (GXID, Global Transaction ID) assignment,
    - Gather transaction status from all the coordinators and maintain snapshot data,
    - Provide snapshot data to each transaction/statement (Global Snapshot).
  - Extract global value providing feature such as
    - Sequence
    - Time/sysdate (future)



# GTM and PG-XC Transaction Management





## GXID and Snapshot

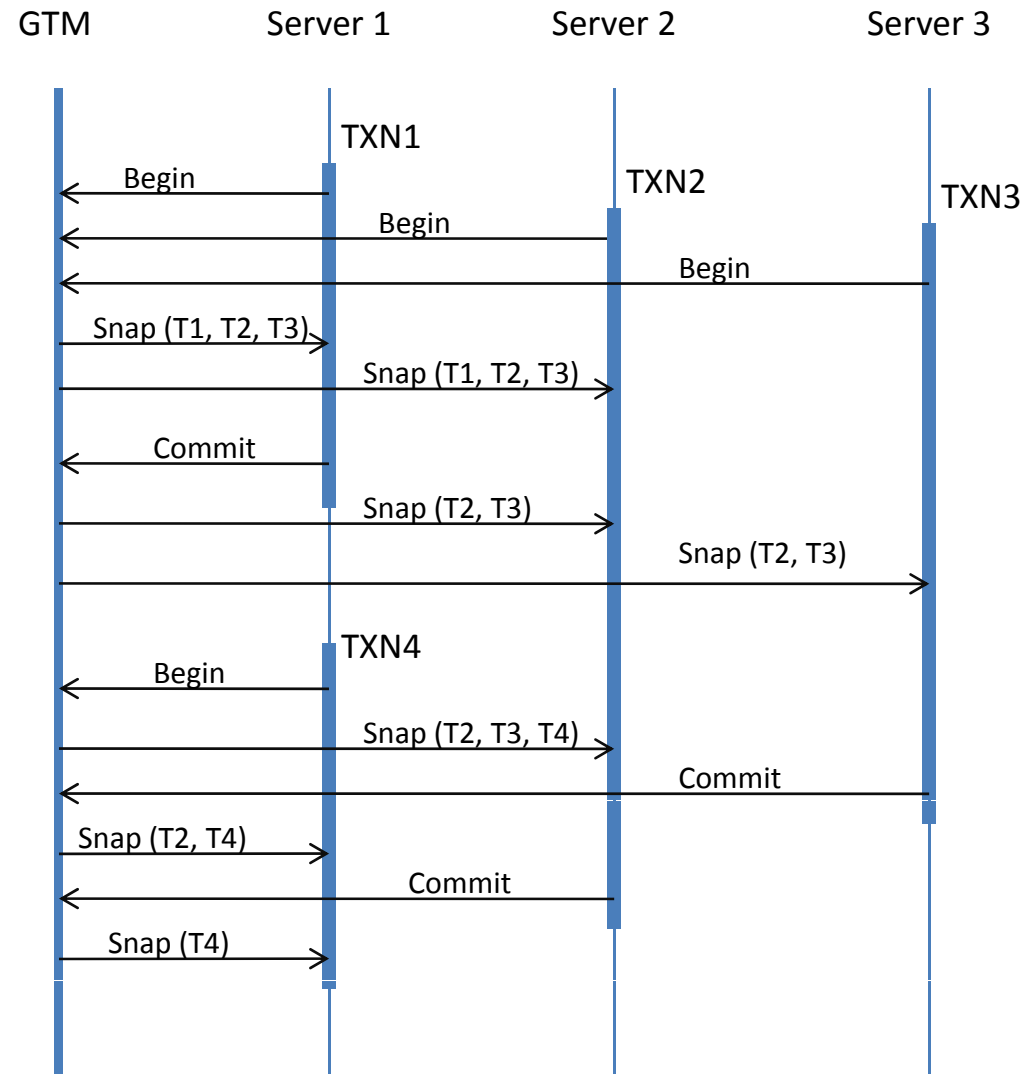
- GXID
  - Unique Transaction ID in the system
- Global Snapshot
  - Includes snapshot information of transactions in other coordinators.



- Data node can handle transactions from different coordinators without consistency problem.
- Visibility is maintained globally, same as standalone PostgreSQL.



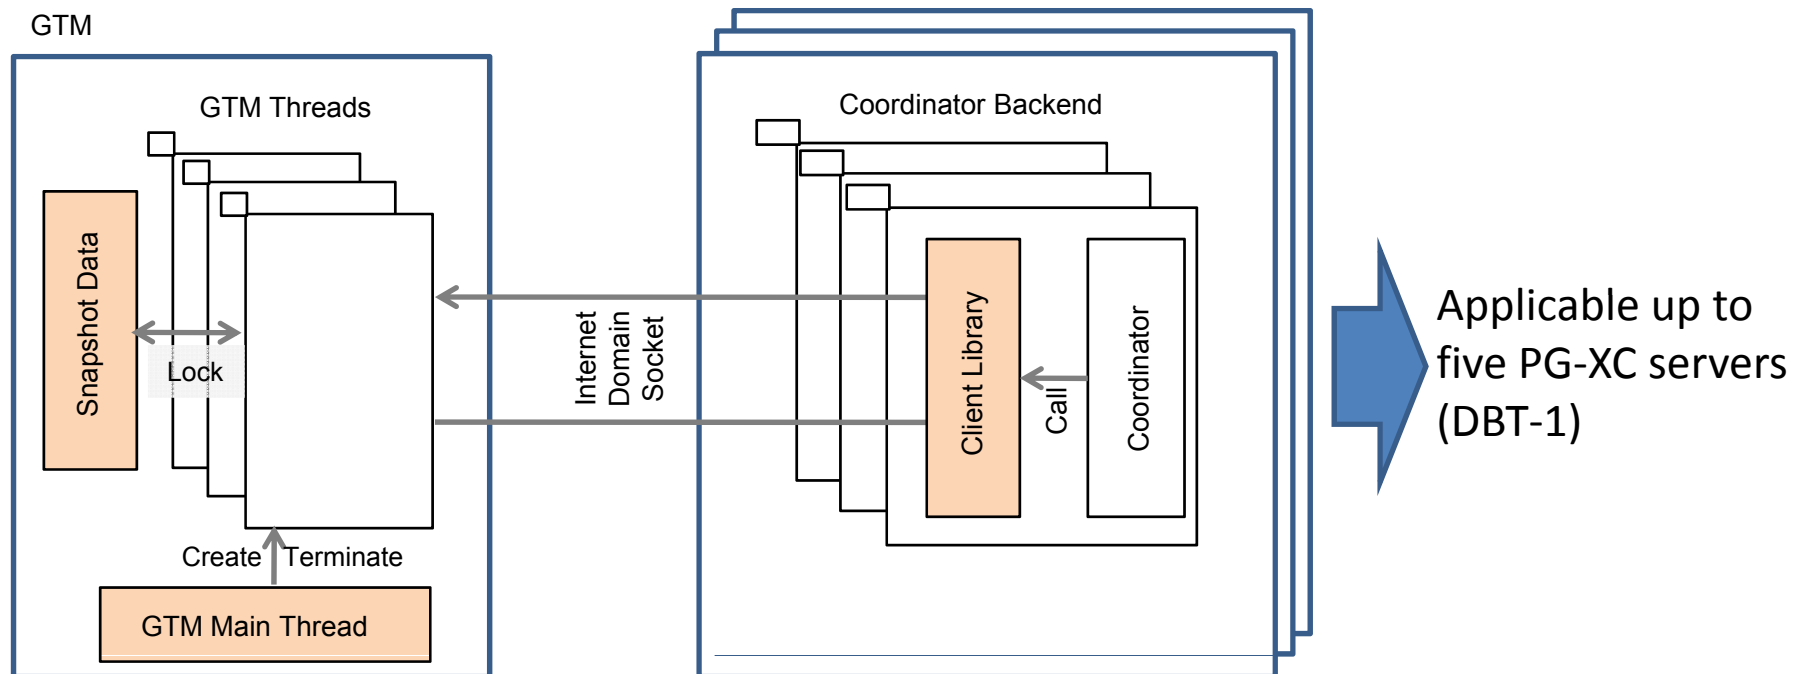
# Outline of PG-XC Transaction Management





# Can GTM be a Performance Bottleneck?

- Depending on implementation
  - Current Implementation

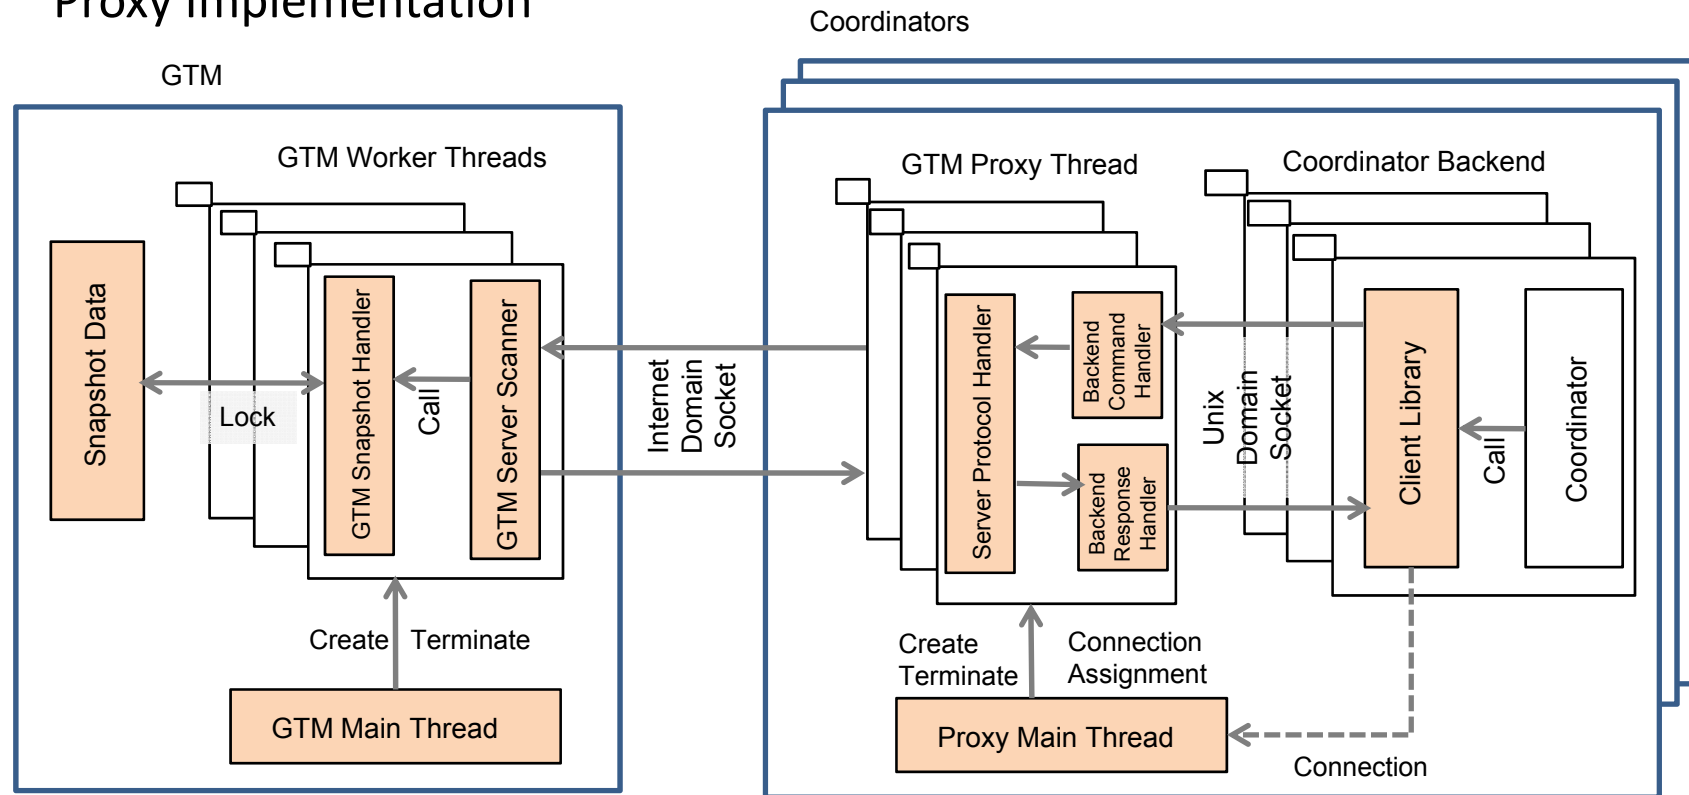


- Large snapshot size and number
- Too many interaction between GTM and Coordinators



# Can GTM be a Performance Bottleneck?

- Proxy Implementation

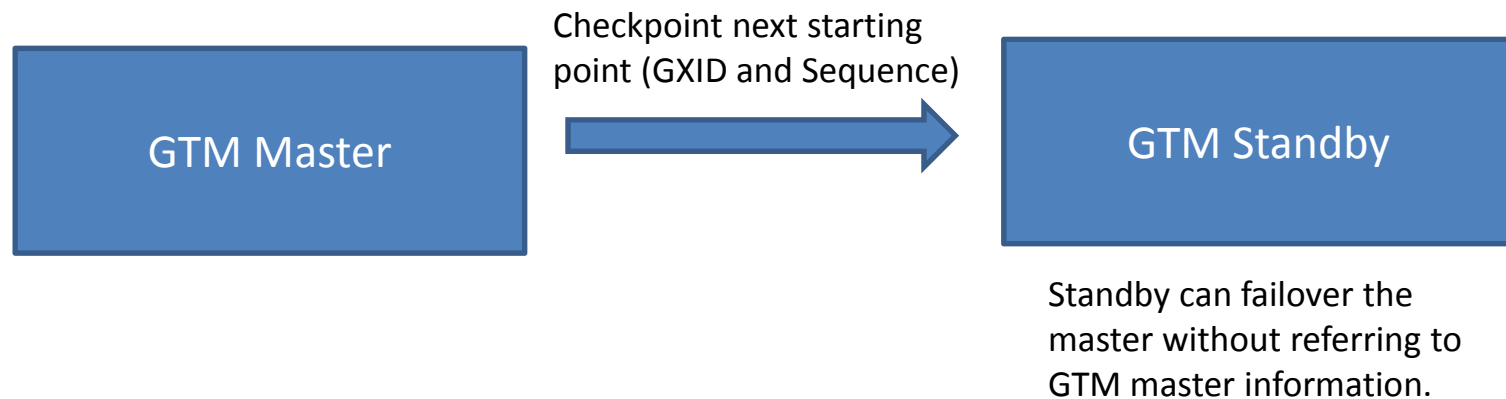


- Very good potential
  - Request/Response grouping
  - Single representative snapshot applied to multiple transactions
- Maybe applicable for more than ten PG-2 servers



## Can GTM be a SPOF?

- Simple to implement GTM standby



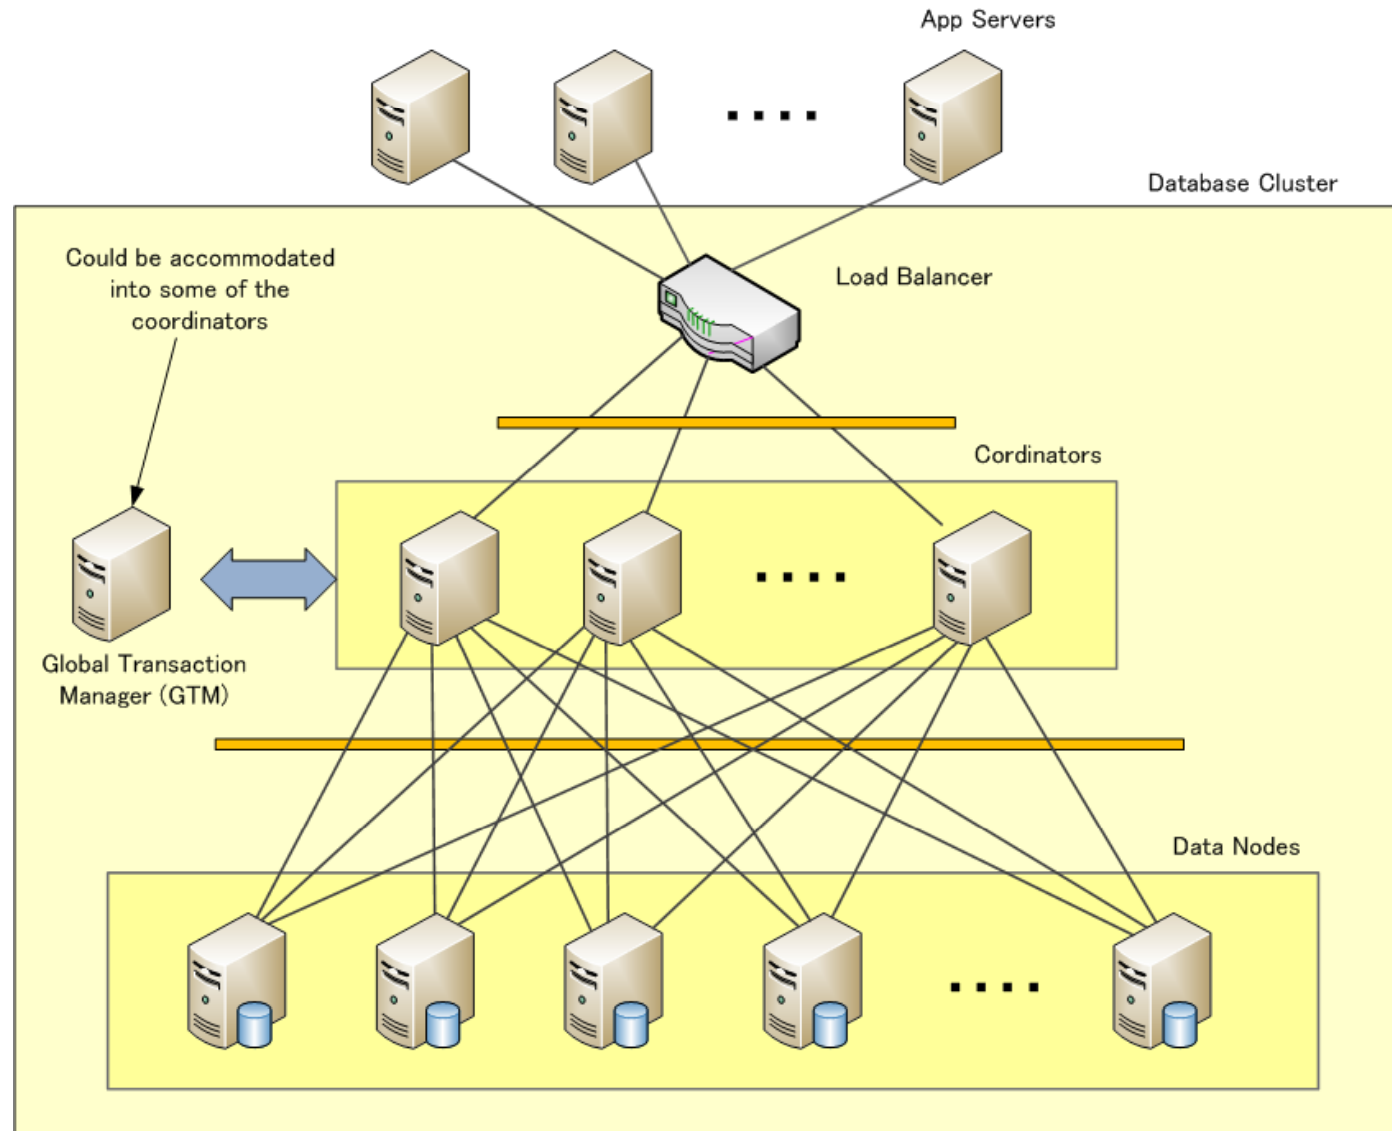


# Coordinator/Data Node Internals





# Reference Architecture





## Coordinator Overview

- Based on PostgreSQL 8.4.3
- Accepts connections from clients
- Parses requests
- Examines requests, reroutes to Data Nodes
- Interacts with Global Transaction Manager
- Uses pooler for Data Node connections
- Sends down XIDs and snapshots to Data Nodes
- Uses two phase commit if necessary



## Data Node Overview

- Based on PostgreSQL 8.4.3
- Where user created data is actually stored
- Coordinators (not clients) connects to Data Nodes
- Accepts XID and snapshots from Coordinator
- The rest is fairly similar to vanilla PostgreSQL



## Postgres-XC Request Handling

- Data Distribution
- Pooler
- Statements
  - Only involve nodes as needed
  - Proxy efficiently
  - If multiple nodes, issue query simultaneously
  - Global MVCC
- Transactions




# Data Distribution

## Distribution Types:

- Hash Partitioning
- Round Robin
- Replicated

Coordinator

Long term: custom, range partitioning





## Connection Pooling

- The Coordinator forks off a pooler process for managing connections to the Data Nodes
- Coordinator obtains connections from pooler process as needed
  - Not every transaction needs all Data Nodes
- At commit time, Coordinator returns connections to the pool
- As we add clients and multiple Coordinators, we want to prevent an explosion of required connections at the data node level by pooling instead




## Statement Handling

- Only basic statements currently handled
  - (no cross-node joins yet)
- Use distribution information in Coordinator
- If more than one Data Node, send down statement to all simultaneously
- Recognize singleton statements
- Recognize single-step statements
- Handle replicated tables
- Use two phase commit
  - (and use only when necessary)



# Statement Handling - Execution







## Queries with Replicated Tables

- Choose a node via round robin to execute on
- Recognize queries with joins between replicated tables

```
SELECT *  
  FROM reptab1 r1 INNER JOIN reptab2 r2  
    ON r1.col1 = r2.col2
```

- For write operations, use all nodes and two phase commit



# Statement Handling - Execution

Table tab1 is hash partitioned on col1

SELECT \*  
FROM tab1  
WHERE col1 = 5

Coordinator

Only use single node  
when possible.

Proxy result efficiently

Data Node

Data Node

Data Node

Data Node



# Statement Handling - Execution

Table tab1 is hash partitioned on col1

SELECT \*  
FROM tab1  
WHERE somecol <>  
10

Coordinator

No condition allows for a single node, send query to all Data Nodes.

Proxies 'D' DataRow messages.

Collects 'C' CommandComplete, when received from all, sends one 'C' message to client.

Data Node

Data Node

Data Node

Data Node



## Queries with Partitioned Tables

- Check WHERE clause to see if we can execute on one node
- Recognize queries with joins with replicated tables

```
SELECT *  
  FROM tab1 t INNER JOIN reftab1 r  
    ON t.col2 = r.col3  
 WHERE t.col1 = 1234
```

- Recognize queries with joins on respective partitioned columns

```
SELECT *  
  FROM tab1 t1 INNER JOIN tab2 t2  
    ON t1.col1 = t2.col1  
 WHERE t.col1 = 1234
```



## Visibility and Data Node Handling

- When the first statement of a transaction needs to execute, a global XID is obtained from GTM
- Each time a new Data Node connection joins a transaction, the Coordinator sends down a GXID to the Data Node
- Each statement execution requires a new snapshot being obtained from GTM
- Before sending down a SQL statement, the Coordinator first passes down a snapshot to the Data Nodes



## Transactions and Data Node Handling

- The Coordinator tracks read and write activity\*
- At commit time
  - If we have only written to one Data Node, we simply issue commit to the node
  - If we have written to more than one Data Node, we use two phase commit
  - If no Data Nodes have been written to, we do not send down any commit

\*Stored functions could theoretically write to DB



# Transaction Handling Considerations

- Distributed transactions and two phase commit (2PC)
- Distributed Multi-Version Concurrency Control
  - Global Snapshots
  - Autovacuum
    - exclude XID in global snapshots
  - ANALYZE
  - Future optimization
  - CLOG
    - Careful when extending, not all transactions are on all nodes



## Aggregate Handling


- Traditional PostgreSQL in Two Phases:
  - Transition Function
  - Finalizer Function
- Postgres-XC uses Three Phases:
  - Transition Function
  - Collector Function
  - Finalizer Function





# Aggregate Handling

## Postgres-XC Aggregate Flow





## Aggregate Handling - AVG

- AVG (Average) needs to sum all elements and divide by the count
- Transition

```
arg1[0] += arg2;  
arg1[1]++;  
return arg1;
```
- Combiner (only in Postgres-XC)

```
arg1[0] += arg2[0];  
arg1[1] += arg2[1];  
return arg1;
```
- Finalizer

```
return arg1[0] / arg1[1];
```

Get the sum of the sums  
and the sum of the counts  
from the Data Nodes



## Looking at Code

- Not (yet) overly invasive in PostgreSQL code
  - 8.4.2 → 8.4.3 merged cleanly
- Existing modules use `#ifdef PGXC` to identify Postgres-XC changes
- `IS_PGXC_COORDINATOR` and `IS_PGXC_DATANODE` easily identifies applicable code
- Advanced Coordinator logic in separate modules



# Evaluation




## Postgres-XC Performance Benchmark

- Based on DBT-1
  - Typical Web-based benchmark
  - We had good experience on this
- Changes from the original
  - Changed ODBC to libpq
    - Put much more workload
  - Added distribution keys
    - Can be automatically generated in the future
  - One table divided into two
    - According to the latest TPC-W specification
    - Matches Postgres-XC characteristics

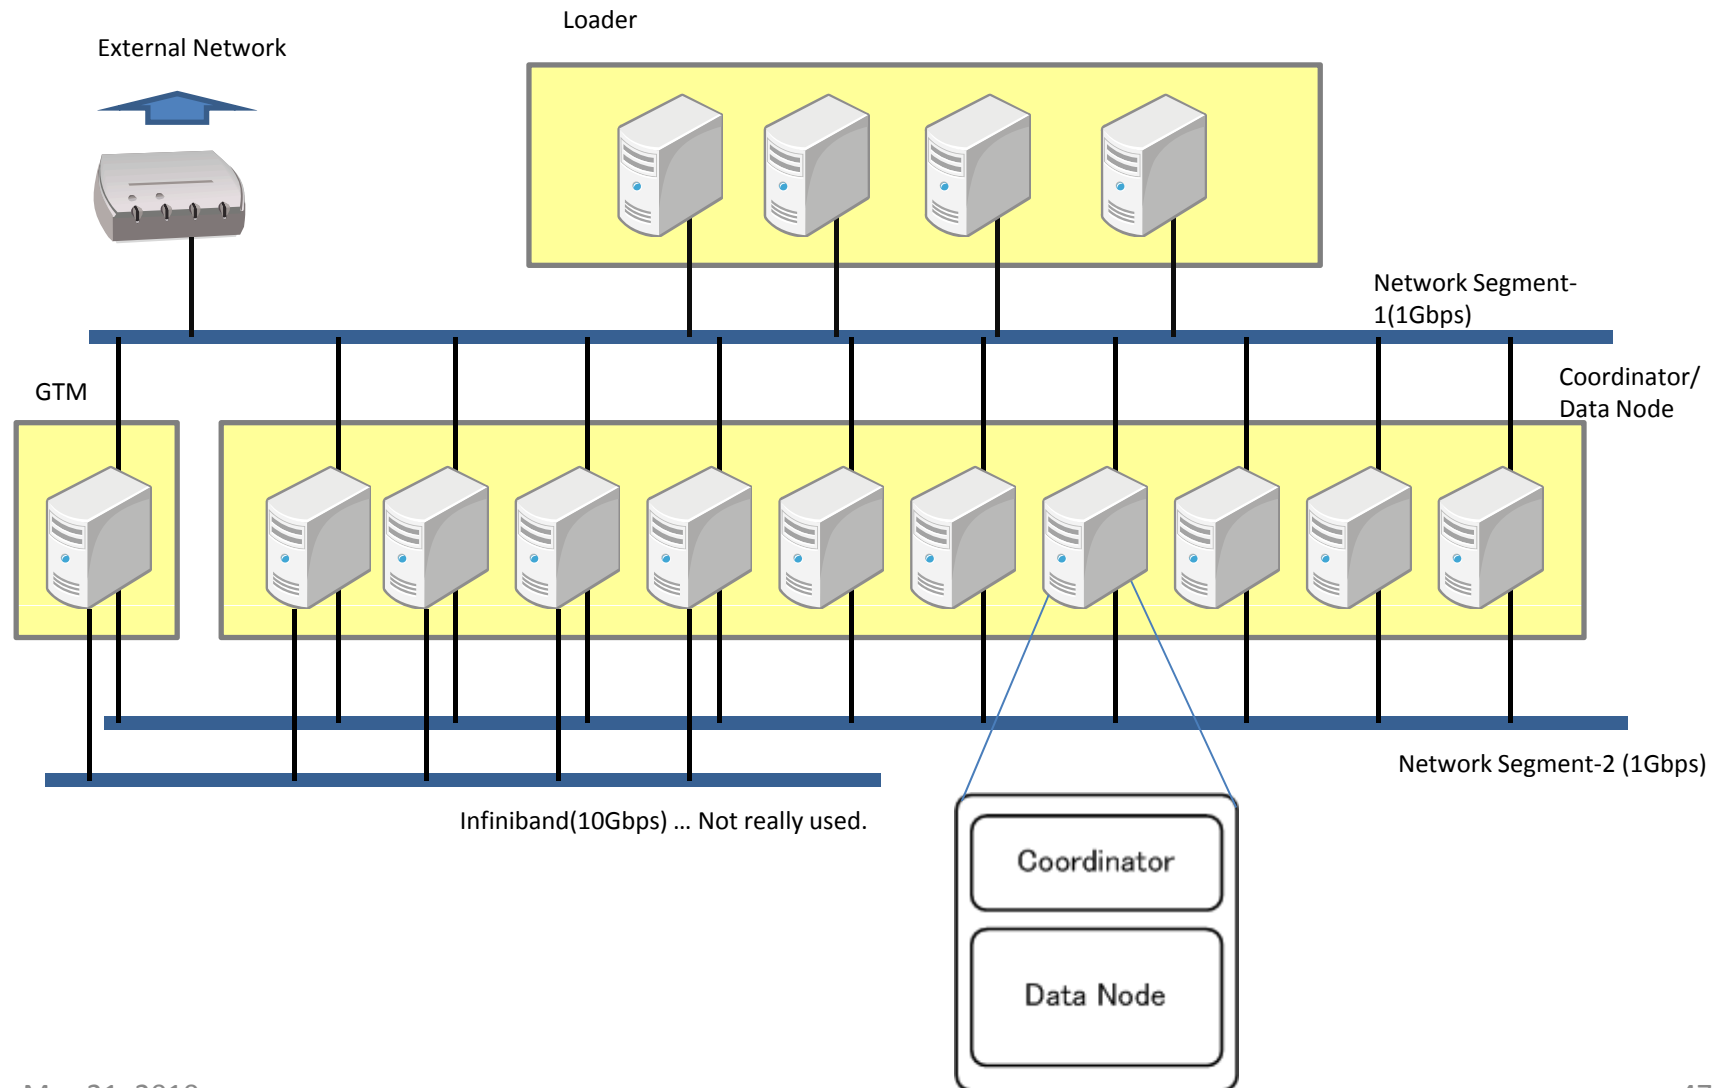


# DBT-1-based Table Structure





# Evaluation Environment





## Server Spec

	Coordinator/Data Node	GTM/Loader
Make	HP Proliant DL360 G6	HP Proliant DL360 G5
CPU	Intel® Xeon® E5504 2.00GHz x 4	Intel® Xeon® X5460 3.16GHz x 4
Cache	4MB	6MB
MEM	12GB	6GB
HDD	146GB SAS 15krpm x 4 ea	146GP SAS 14krpm x 2 ea





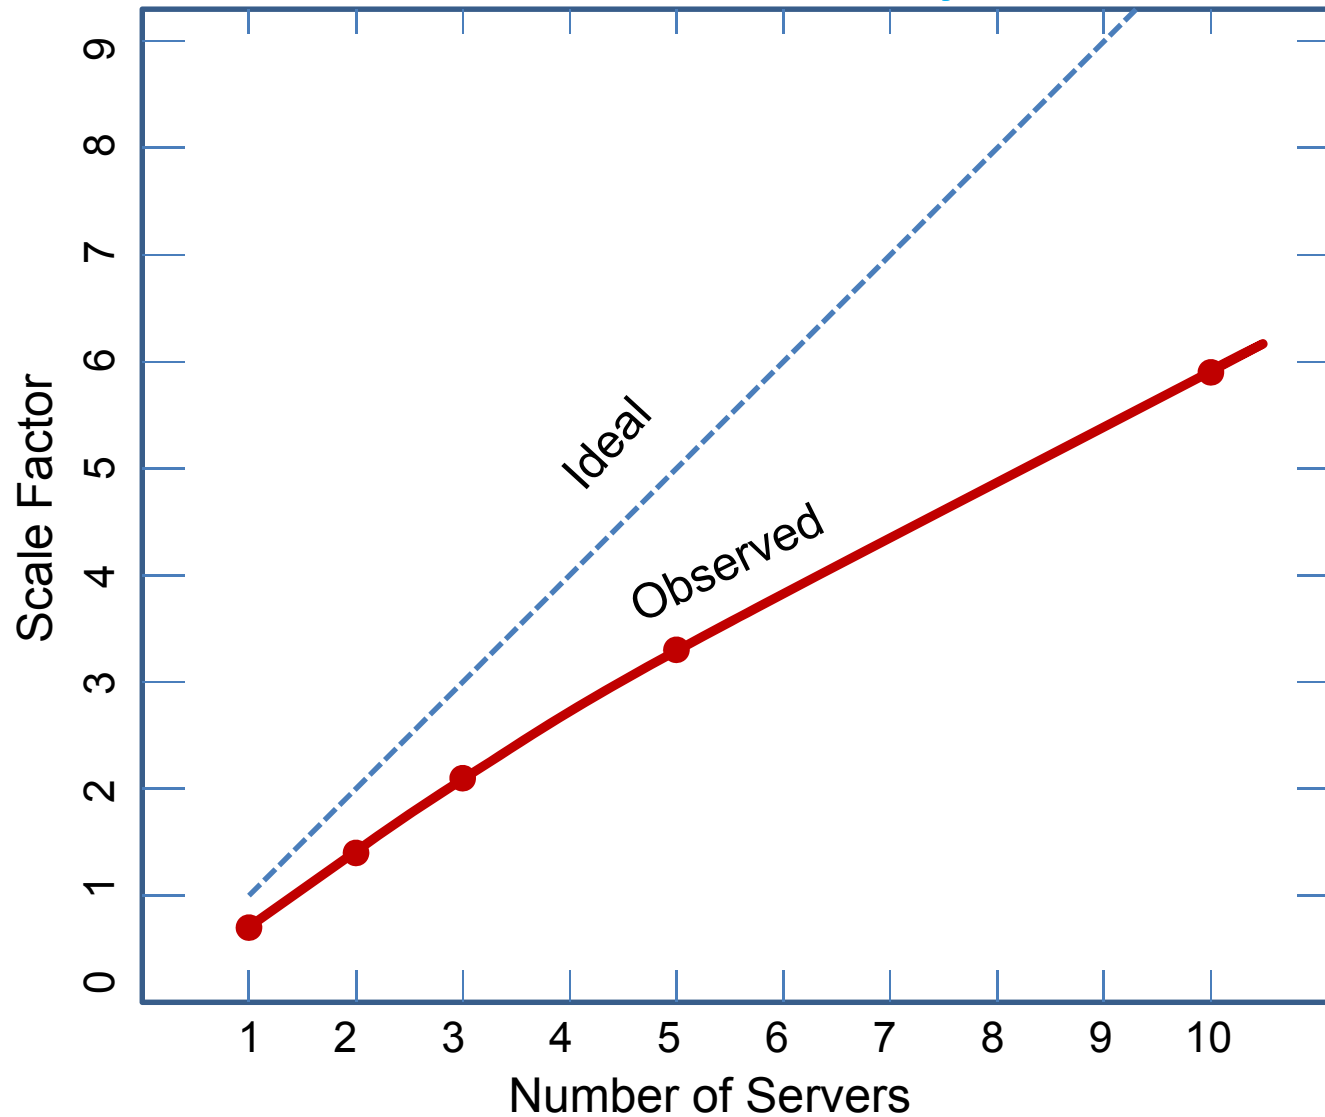
# Evaluation Summary

## Full Load Throughput

Database	Num. of Servers	Throughput (TPS)	Scale Factor
PostgreSQL	1	2,617	1.0
Postgres-XC	1	1,869	0.71
Postgres-XC	2	3,646	1.39
Postgres-XC	3	5,379	2.06
Postgres-XC	5	8,473	3.24
Postgres-XC	10	15,380	5.88

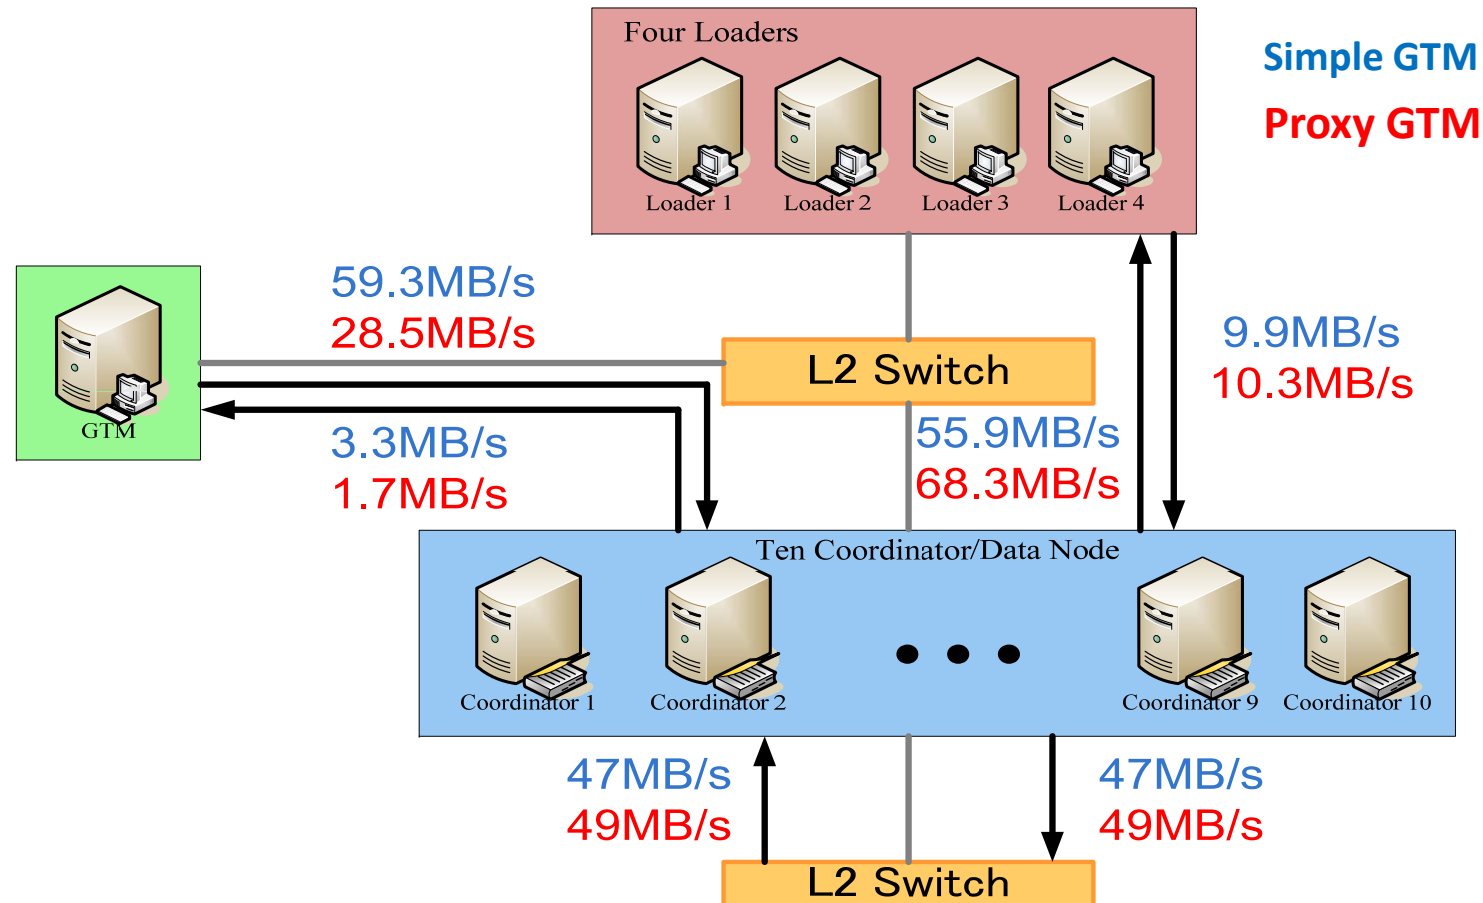


## Scale Factor Summary






# Network Workload





# One Week Test



Vaccum Analyze may become long transactions to affect the throughput.

Reasonably stable in a long run (90% workload)



## Avoiding Long Transactions

- Vacuum
  - Needs GXID
  - Vacuum's GXID need not to appear in local or global snapshot
- Vacuum Analyze
  - Needs GXID
  - GXID should appear in local snapshot
  - GXID need not appear in global snapshot

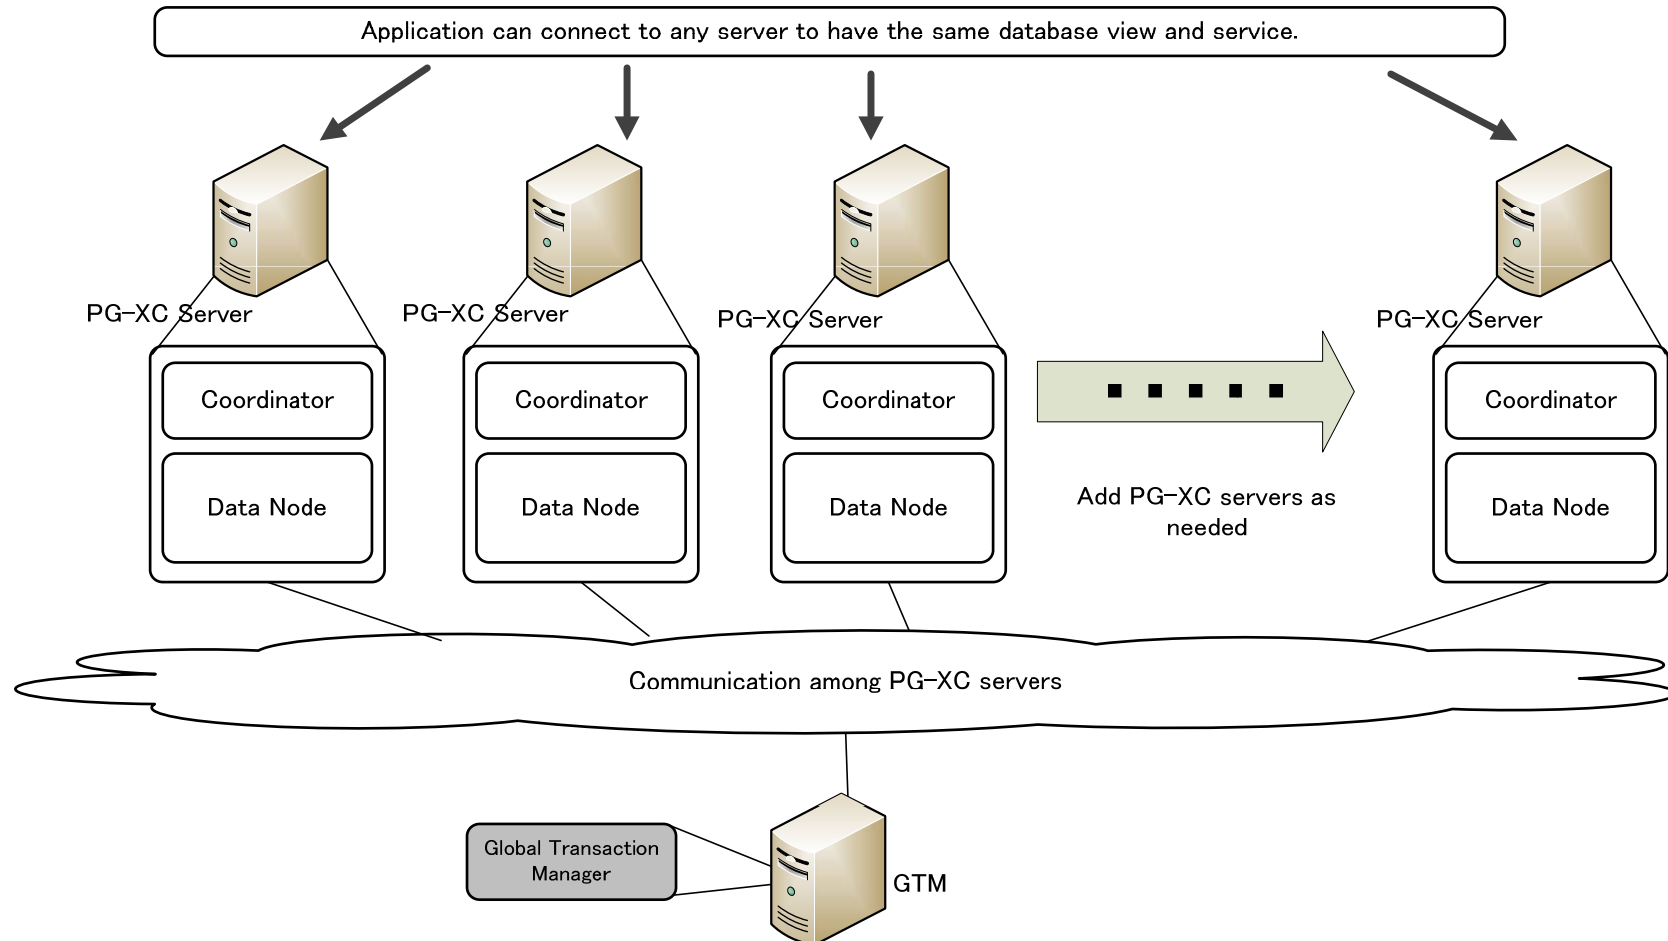


## Evaluation Summary

- PG-XC is reasonably scalable in both read/write.
- Need some tweak to stabilize the performance.
- Network workload is reasonable.
  - GTM Proxy works well.
  - More work is needed to accommodate more servers (thirty or more)
- Fundamentals are established
  - Will continue to extend statement support



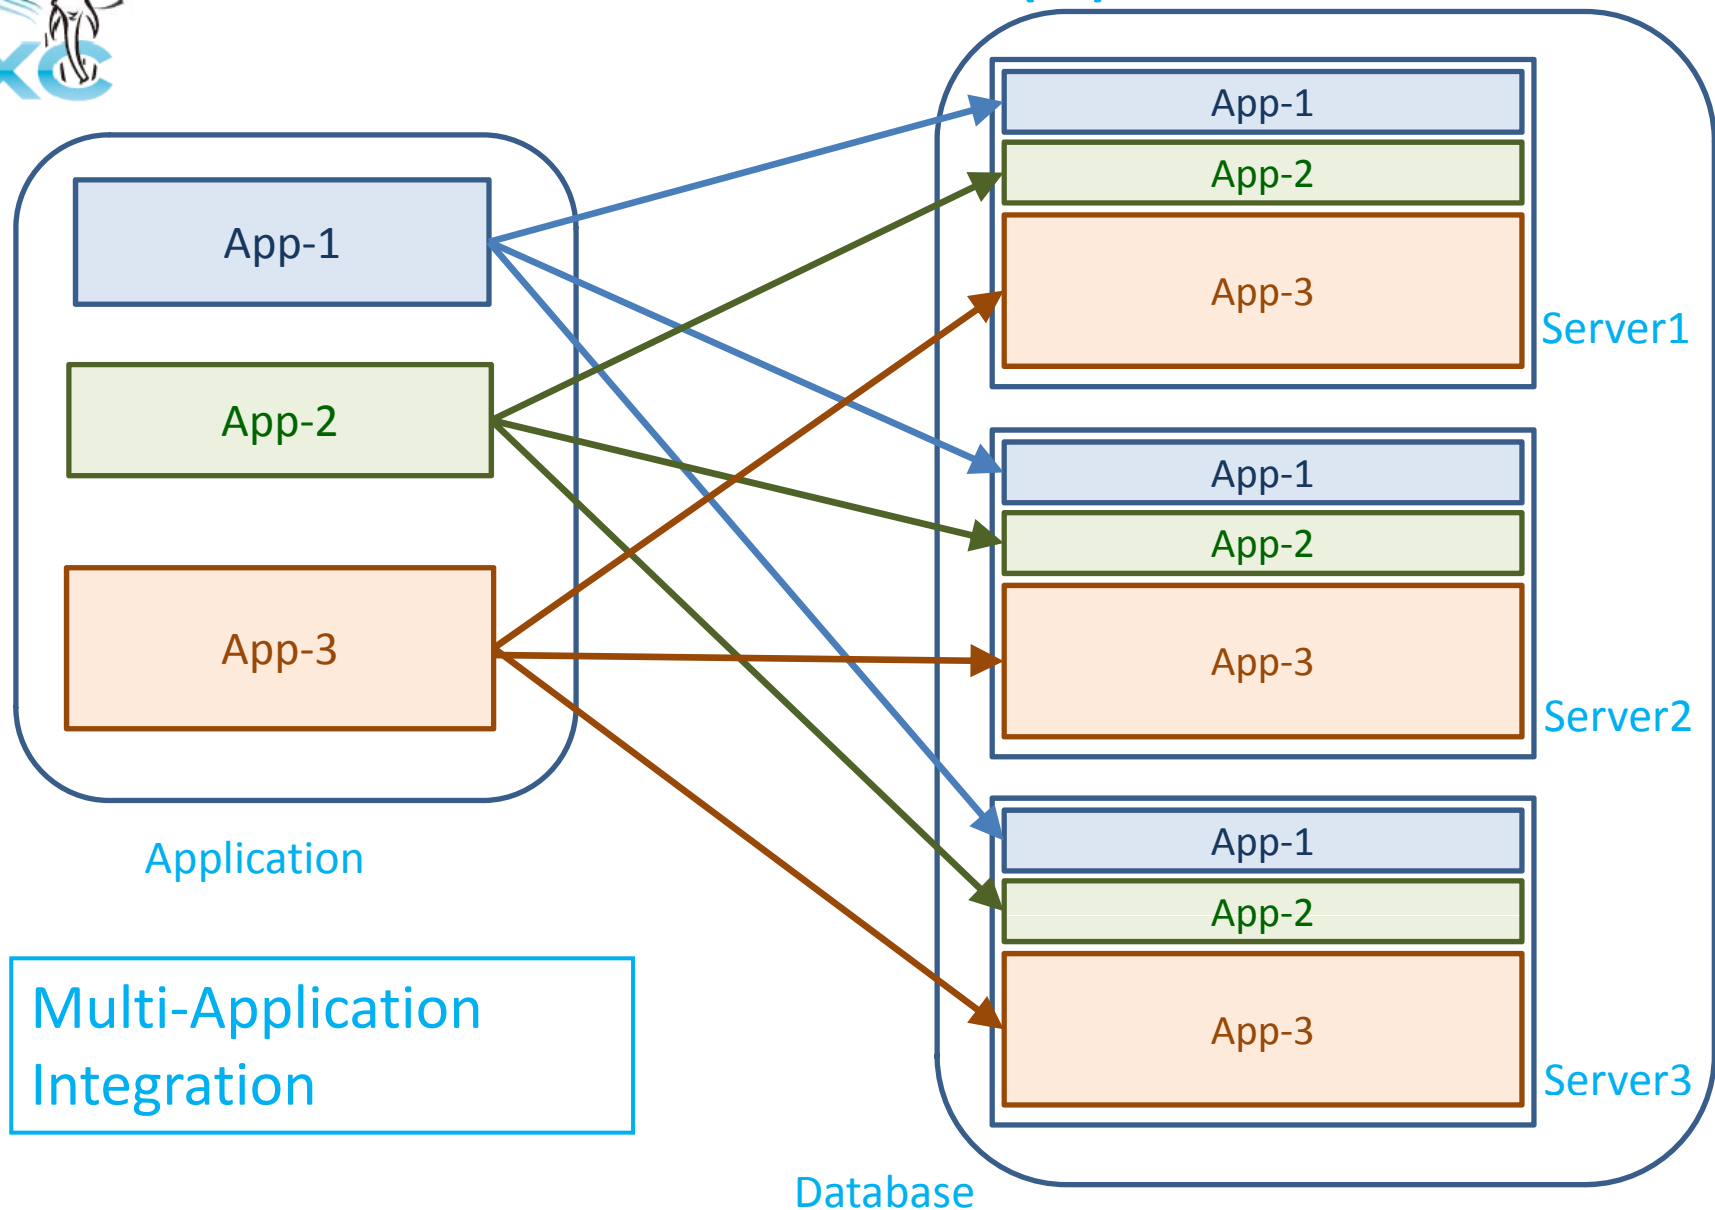
# Possible Use Case (1)



Large Scale Application



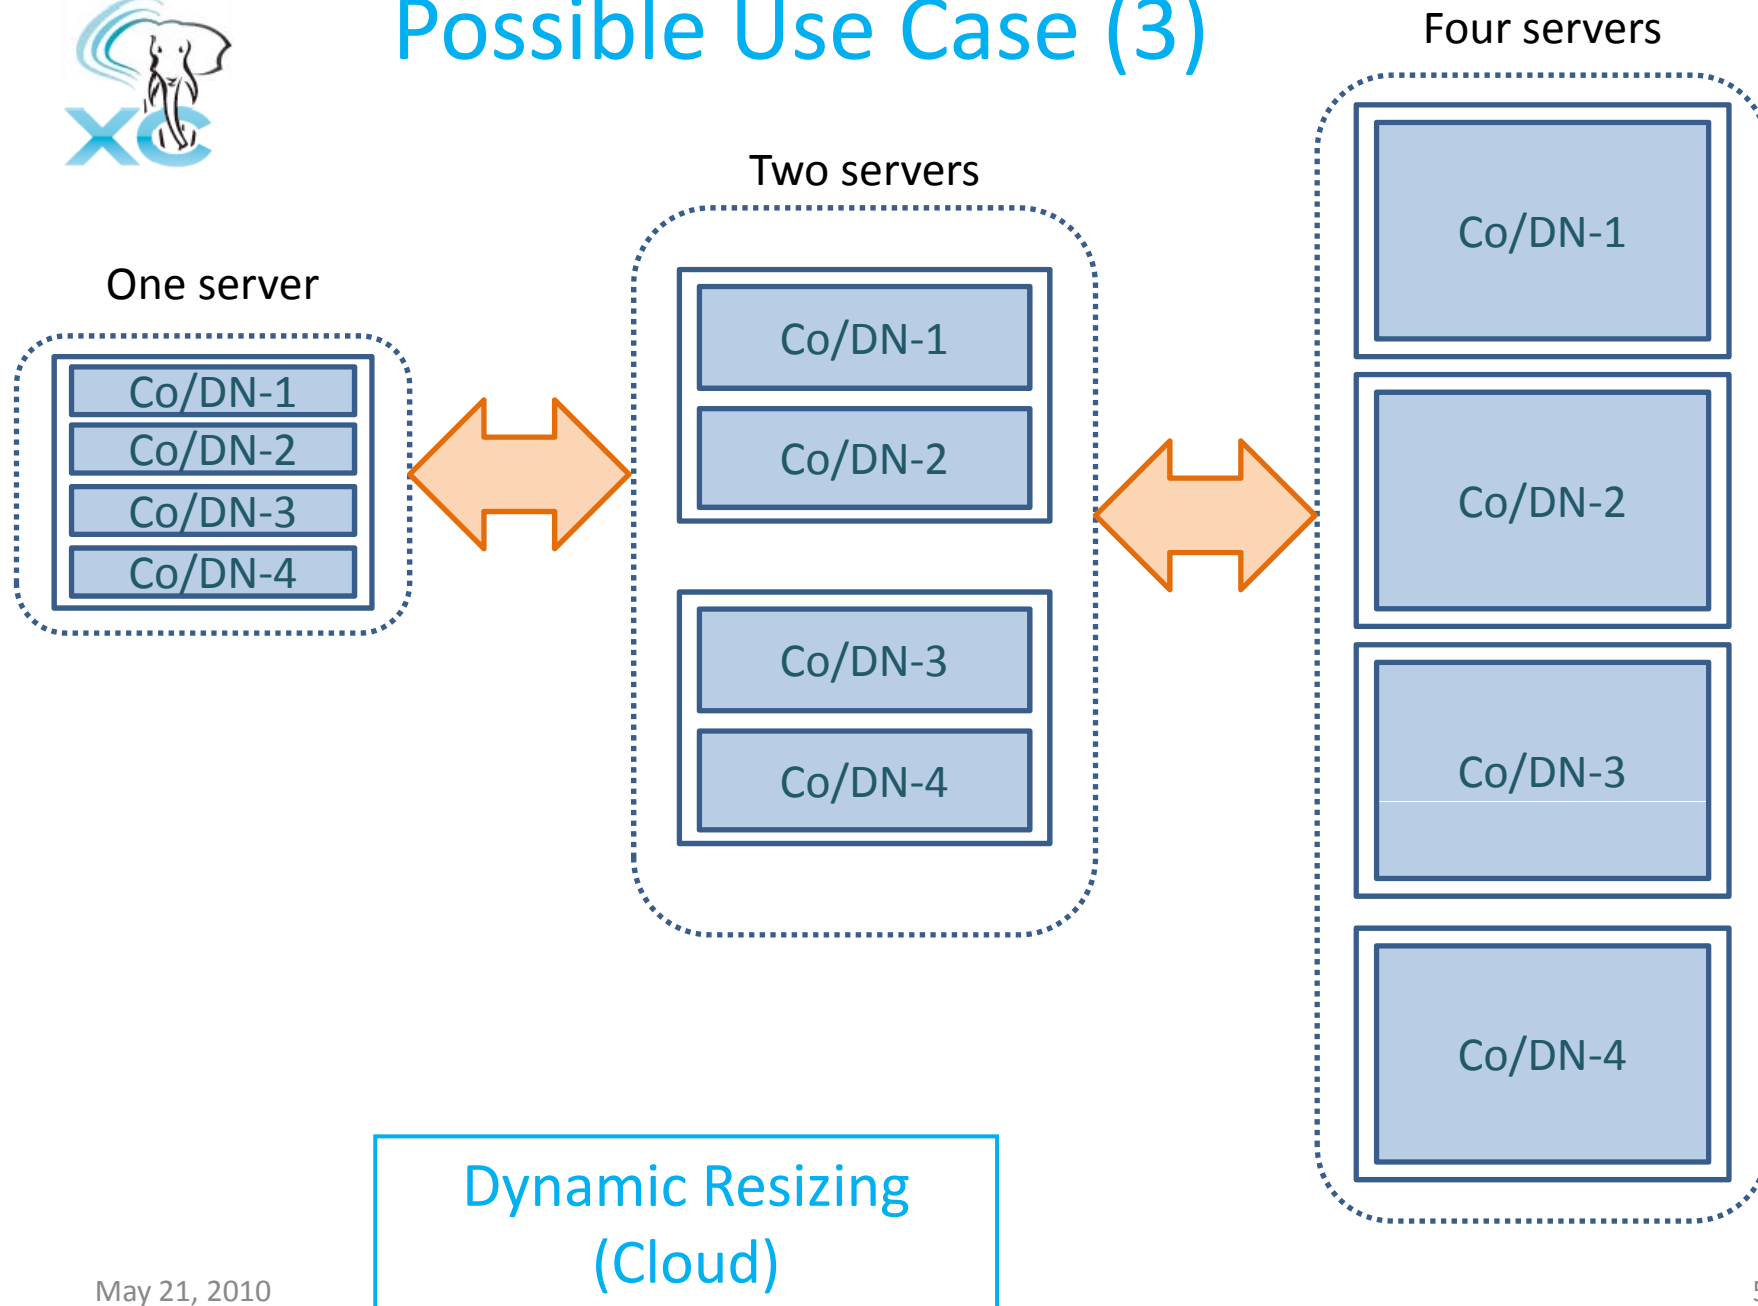
## Possible Use Case (2)







## Possible Use Case (3)





WIP

- WIP up to V.1.0
  - ORDER BY/DISTINCT
  - Stored Functions
  - Subqueries
  - Views, Rules



## Roadmap and Plan (1)

- Toward V.1.1 (Sept., 2010)
  - Cluster-Wide Installer
  - Cluster-Wide Operation Utilities
  - Regression Tests
  - Logical Backup/Restore
  - Basic Cross-Node Operation
  - TEMP Table
  - Extended Query Protocol
    - JDBC
  - Global Timestamp
  - Drivers
    - ECPG, JDBC, PHP, etc.
  - Forward Cursor (w/o ORDER BY)



## Roadmap and Plan (2)

- Beyond V.1.1
  - Physical Backup/Restore
    - PITR
  - Cross-node operation optimization
    - Tuple transfer Infrastructure from node to node
  - More variety of statements
    - INSERT ... FROM SELECT
  - Prepared Statement
  - General Aggregate Functions
  - General Functions
  - Savepoint
  - Session Parameters
  - 2PC from Apps
  - Forward Cursor with ORDER BY
  - Backward Cursor
  - Batch, Statement pushdown
  - Catalog Synchronize with DDL
  - Trigger
  - Global constraints
  - Tuple relocation
    - Distribute key update



## Interesting Remarks with SQL/MED

- Postgres-XC vs SQL/MED
  - Tightly-coupled vs Autonomous
  - R/W vs Read-centric
  - Single application vs Independent applications

## Nevertheless...

- Postgres-XC and SQL/MED shares cross-node operation
  - First SQL/MED effort is applicable to Postgres-XC as well.
    - Need to add global transaction feature.
    - Targeted to V.1.1.
  - Upcoming Postgres-XC effort can be brought to SQL/MED.
    - Postgres-XC targets only (a kind of) PostgreSQL database so our SQL/MED may need some more work to apply.



## Developers Welcome

- We welcome people helping the project.
  - Each issue in WIP and the roadmap is composed of small manageable pieces.
  - If you are interested in the project, please contact us.
- Project Home Page
  - <http://postgres-xc.sourceforge.net/>
- Contact
  - [koichi.szk@gmail.com](mailto:koichi.szk@gmail.com)
  - [mason.sharp@gmail.com](mailto:mason.sharp@gmail.com)



Thank You Very Much